

PIOTR WOŚ

SHARED MEMORY LINKED LIST

<https://github.com/ptr97/shared-memory-linked-list>

PROJECT DESCRIPTION

- ▶ **"Complex objects with pointers (say linked list) in shared memory"**
- ▶ For simplicity let's say that:
 - ▶ we have only one producer and only one consumer
 - ▶ producer and consumer both work sequentially - firstly producer adds something to the list, then consumer reads from there

PROJECT DESCRIPTION

- ▶ I will show you solution which uses arithmetic on offsets from the beginning of shared memory
- ▶ And if we still have some time left I will tell you about hardest one which should work (and almost works ;)) with use of arithmetic on pointers

PROJECT DESCRIPTION

- ▶ Shared memory
 - ▶ mmap
 - ▶ shmget, shmat
 - ▶ I chosen mmap because of easy, natural access to it (you can easily look inside the file)
- ▶ Why structures like linked list aren't easy to store in shared memory?

THE CHALLENGE

- ▶ One and only one challenge is the fact that virtual address of shared memory may be and almost always is different each time.
- ▶ Is it all?
- ▶ Yea... pointer arithmetic...
- ▶ And... many ideas for solution.

THE SOLUTION(S)



```
1  #include <iostream>
2  #include "SharedMemory.h"
3  #include "List.h"
4
5
6  int main()
7  {
8      const uint shmBlockSize = 4096 * 2;
9
10     List<int> list = List<int>::createListInShm("database.db", shmBlockSize);
11     list.add(10);
12     list.add(20);
13     list.add(30);
14     list.add(40);
15     list.print();
16     std::cout << "list.exists(20) == " << list.exists(20) << std::endl;
17     list.remove(20);
18     list.remove(10);
19     list.print();
20     shmBlock::freeShm(shmBlockSize);
21
22     std::cout << std::endl << "Now we will read linked list from memory" << std::endl;
23     List<int> listFromMem = List<int>::readListFromMemory("database.db");
24     listFromMem.print();
25     shmBlock::freeShm(shmBlockSize);
26
27     return 0;
28 }
```

```
Memory block of size 8192 allocated at: 0x109533000
List capacity: 682
freeOffset = 0
freeOffset = 1
freeOffset = 2
freeOffset = 3
Current: 40
Current: 30
Current: 20
Current: 10
list.exists(20) == 1
removing node with value 20
removing node with value 10
Current: 40
Current: 30
Memory block has been released...
```



```
Now we will read linked list from memory
Memory block of size 8192 read at: 0x109533000
List capacity: 682
Current: 40
Current: 30
Memory block has been released...
```

SHARED MEMORY LINKED LIST

THE SOLUTION(S)



```
Memory block of size 8192 allocated at: 0x102f51000
List capacity: 682
freeOffset = 0
freeOffset = 1
freeOffset = 2
freeOffset = 3
Current: 40
Current: 30
Current: 20
Current: 10
Memory block has been released...
```



```
Memory block of size 8192 read at: 0x10dd85000
List capacity: 682
Current: 40
Current: 30
Current: 20
Current: 10
Memory block has been released...
```

SHARED MEMORY LINKED LIST

THE SOLUTION(S)

```
1  #ifndef __SHARED_MEMORY_H__
2  #define __SHARED_MEMORY_H__
3
4  #include <iostream>
5  #include <cstring>
6  #include <cstdlib>
7  #include <sys/types.h>
8  #include <sys/stat.h>
9  #include <sys/fcntl.h>
10 #include <sys/mman.h>
11 #include <unistd.h>
12
13
14 class shmBlock
15 {
16 public:
17     static char * startPtr;
18
19     static void allocateMemory(const char * key, uint size);
20     static uint readFromMemory(const char * key);
21     static void freeShm(uint size);
22 };
23
24 #endif
```

```
void shmBlock::freeShm(uint size) {
    if(munmap(shmBlock::startPtr, static_cast<off_t>(size))) {
        std::cout << "unmap error" << std::endl;
        exit(-1);
    }
    std::cout << "Memory block has been released..." << std::endl;
}
```

```
1  #include "SharedMemory.h"
2
3  char * shmBlock::startPtr;
4
5  void shmBlock::allocateMemory(const char * key, uint size) {
6      const int fd = open(key, O_RDWR | O_CREAT | O_TRUNC, S_IRUSR | S_IWUSR);
7      if( fd < 0 ) {
8          std::cout << "Open DB error" << std::endl;
9          exit(-1);
10     }
11
12     const int status = ftruncate(fd, static_cast<off_t>(size));
13     if(status != 0) {
14         std::cout << "ftruncate error" << std::endl;
15         exit(-1);
16     }
17
18     shmBlock::startPtr = (char *) mmap(0, size, PROT_WRITE, MAP_SHARED, fd, 0);
19     if(shmBlock::startPtr == MAP_FAILED) {
20         std::cout << "mmap failed" << std::endl;
21         exit(-1);
22     }
23
24     memset(shmBlock::startPtr, '\0', static_cast<off_t>(size));
25     std::cout << "Memory block of size " << size << " allocated at: " << (void *) shmBlock::startPtr << std::endl;
26 }
27
28 uint shmBlock::readFromMemory(const char * key) {
29     const int fd = open(key, O_RDONLY, S_IRUSR | S_IWUSR);
30     if(fd < 0) {
31         std::cout << "Open DB error" << std::endl;
32         exit(-1);
33     }
34
35     struct stat file_statistics;
36     fstat(fd, &file_statistics);
37
38     shmBlock::startPtr = (char *) mmap(0, file_statistics.st_size, PROT_READ, MAP_PRIVATE, fd, 0);
39     if(shmBlock::startPtr == MAP_FAILED) {
40         std::cout << "mmap failed" << std::endl;
41         exit(-1);
42     }
43     std::cout << "Memory block of size " << file_statistics.st_size << " read at: " << (void *) shmBlock::startPtr << std::endl;
44     return file_statistics.st_size;
45 }
```


THE SOLUTION(S)

```
16  template <typename T>
17  class List {
18      template <typename U>
19      struct Node {
20          Node(const U u, int offset) {
21              this->value = u;
22              this->nextOffset = offset;
23          }
24
25          U value;
26          int nextOffset = -1;
27          bool saved = true;
28      };
29
30      struct Meta {
31          unsigned int head;
32      };
33  }
```

```
34  public:
35      static List<T> createListInShm(const char * shmName, uint shmBlockSize) {
36          shmBlock::allocateMemory(shmName, shmBlockSize);
37          return List(true, shmBlockSize);
38      }
39
40      static List<T> readListFromMemory(const char * shmName) {
41          uint shmSize = shmBlock::readFromMemory(shmName);
42          return List(false, shmSize);
43      }
```

```
102 private:
103     List(bool newList, uint shmBlockSize) {
104         struct List<T>::Meta & meta = getMeta();
105         setCapacity(shmBlockSize);
106         if (newList) {
107             meta.head = -1;
108             for(int i = 0; i < maxSize; ++i) {
109                 getNode(i).saved = false;
110             }
111         }
112     }
```

THE SOLUTION(S)

```
114 Node<T> & getNode(int offset) {
115     return *reinterpret_cast<Node<T>*>(shmBlock::startPtr + sizeof(Meta) + sizeof(Node<T>) * offset);
116 }
117
118 Meta & getMeta() {
119     return *reinterpret_cast<Meta *>(shmBlock::startPtr);
120 }
121
122 int findFreeOffset() {
123     int next_offset = 0;
124     while (getNode(next_offset).saved != false) {
125         if (next_offset > maxSize) {
126             std::cout << "ERROR: Out of shm memory" << std::endl;
127             exit(-1);
128         }
129         ++next_offset;
130     }
131     return next_offset;
132 }
133
134 void setCapacity(uint shmBlockSize) {
135     maxSize = (shmBlockSize - sizeof(Meta)) / sizeof(Node<T>);
136     std::cout << "List capacity: " << maxSize << std::endl;
137 }
138
139 int maxSize;
140 };
```

THE SOLUTION(S)

```
45 void add(T item){
46     int freeOffset = findFreeOffset();
47
48     std::cout << "freeOffset = " << freeOffset << std::endl;
49     Meta & meta = getMeta();
50     Node<T> & newNode = getNode(freeOffset);
51     newNode.value = item;
52     newNode.nextOffset = meta.head;
53     newNode.saved = true;
54
55     meta.head = freeOffset;
56 }
57
58 bool exists(T item) {
59     Meta & meta = getMeta();
60     int current = meta.head;
61     while (current != -1) {
62         const Node<T> & currentNode = getNode(current);
63         if (currentNode.value == item) {
64             return true;
65         } else {
66             current = currentNode.nextOffset;
67         }
68     }
69     return false;
70 }
```

```
72 bool remove(T item) {
73     Meta & meta = getMeta();
74     int current = meta.head;
75     int last = meta.head;
76     while (current != -1) {
77         Node<T> & currentNode = getNode(current);
78         if (currentNode.value == item) {
79             std::cout << "removing node with value " << currentNode.value <<
80             currentNode.saved = false;
81             Node<T> & lastNode = getNode(last);
82             lastNode.nextOffset = currentNode.nextOffset;
83             return true;
84         } else {
85             last = current;
86             current = currentNode.nextOffset;
87         }
88     }
89     return false;
90 }
91
92 void print() {
93     Meta & meta = getMeta();
94     int current = meta.head;
95     while (current != -1) {
96         const Node<T> & currentNode = getNode(current);
97         std::cout << "Current: " << currentNode.value << std::endl;
98         current = currentNode.nextOffset;
99     }
100 }
```

SHARED MEMORY LINKED LIST

RESULT

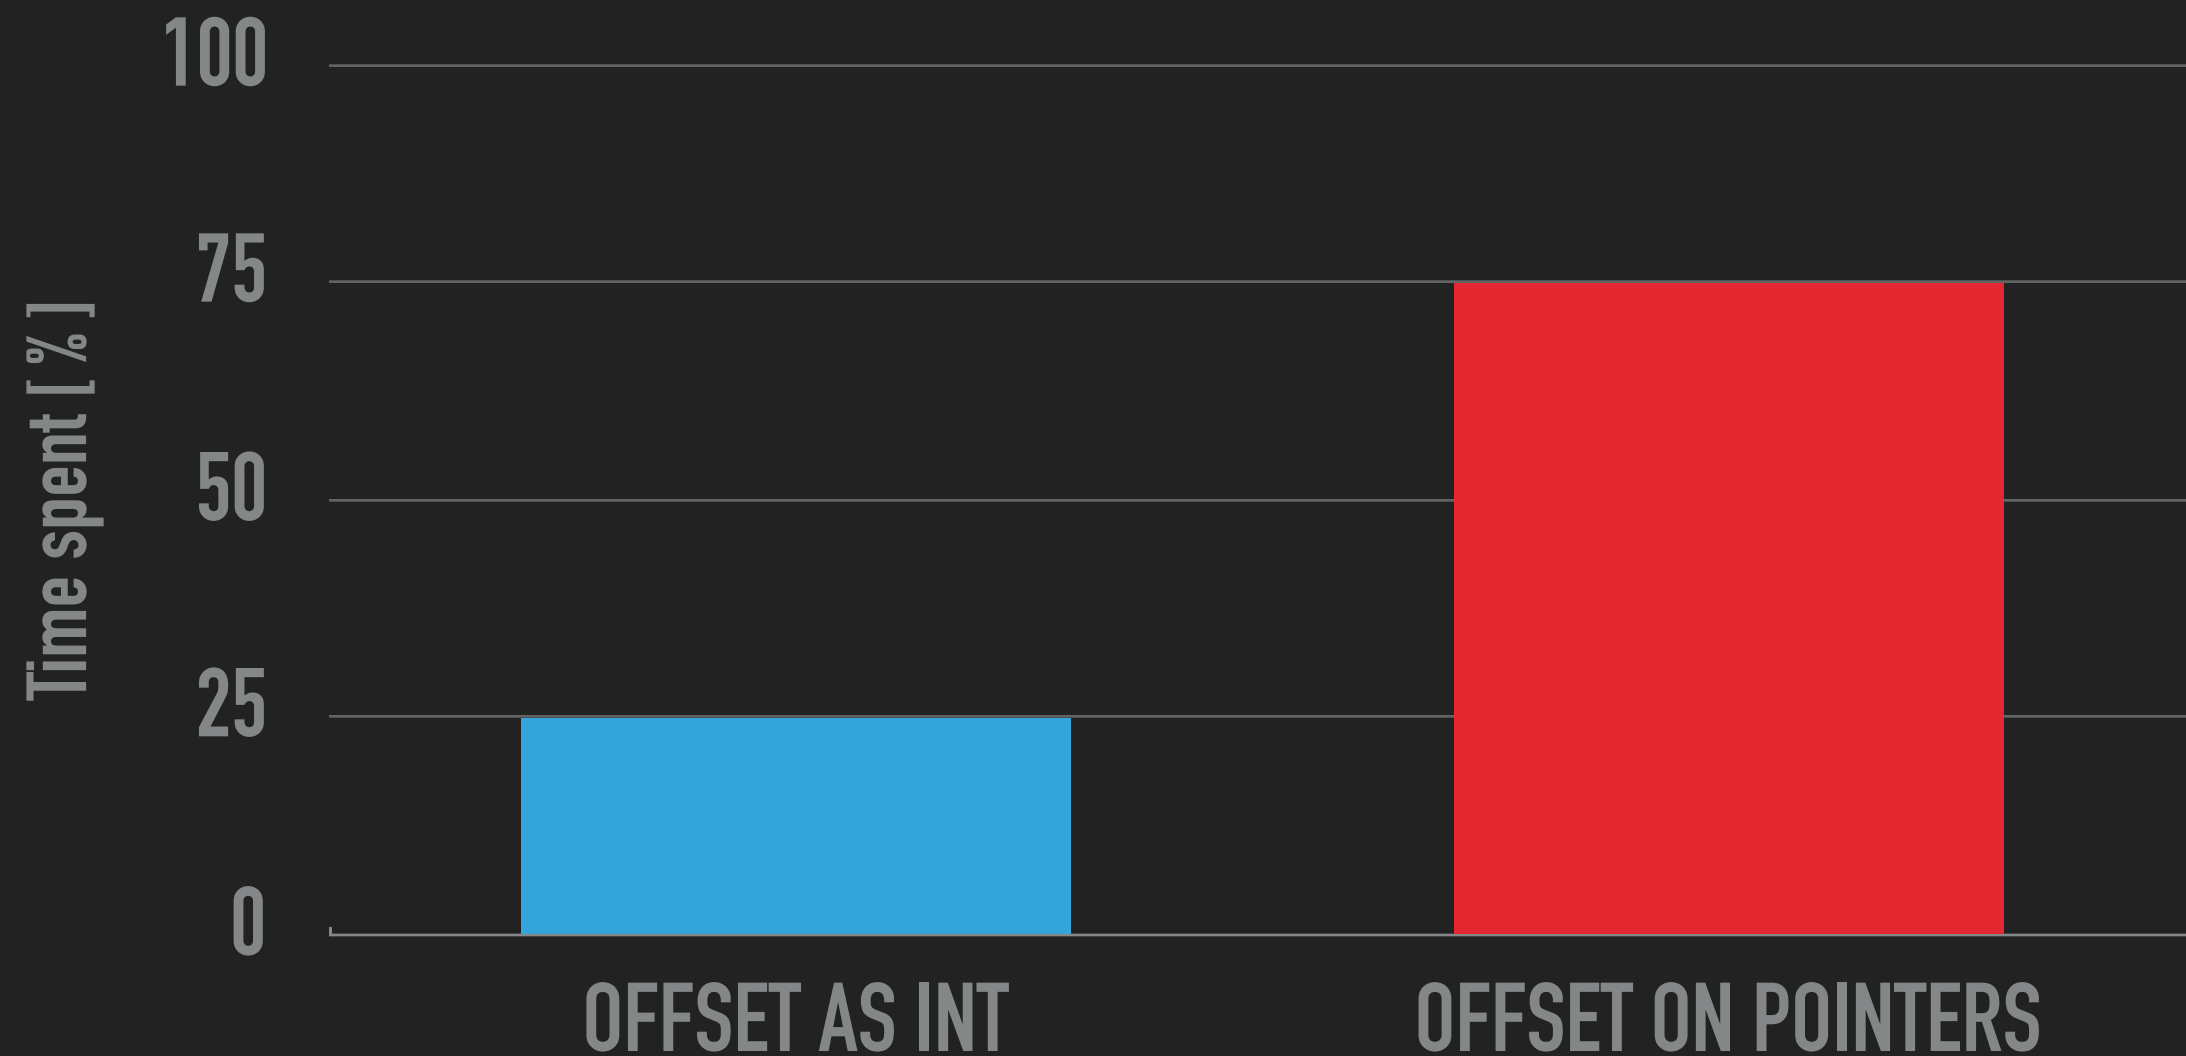
```
list.add(0);
list.add(10);
list.add(20);
list.add(30);
list.add(40);
list.add(50);
list.add(60);
list.add(70);
list.print();
std::cout << "list.exists(20) == " << list.exists(20) << std::endl;
list.remove(40);
list.remove(50);
list.add(150);
list.add(140);
list.remove(50);
list.add(160);
list.print();
shmBlock::freeShm(shmBlockSize);
```

```
Memory block of size 8192 allocated at: 0x10ced2000
List capacity: 682
freeOffset = 0
freeOffset = 1
freeOffset = 2
freeOffset = 3
freeOffset = 4
freeOffset = 5
freeOffset = 6
freeOffset = 7
Current: 70
Current: 60
Current: 50
Current: 40
Current: 30
Current: 20
Current: 10
Current: 0
list.exists(20) == 1
removing node with value 40
removing node with value 50
freeOffset = 4
freeOffset = 5
freeOffset = 8
Current: 160
Current: 140
Current: 150
Current: 70
Current: 60
Current: 30
Current: 20
Current: 10
Current: 0
Memory block has been released...
```

```
Now we will read linked list from memory
Memory block of size 8192 read at: 0x10ced2000
List capacity: 682
Current: 160
Current: 140
Current: 150
Current: 70
Current: 60
Current: 30
Current: 20
Current: 10
Current: 0
Memory block has been released...
```

HARDEST ISSUE

- ▶ As I said at the beginning pointer arithmetic



SHARED MEMORY LINKED LIST

**THANKS FOR
LISTENING**